"Automated Flight Delay Processing Pipeline" project using **AWS Step Functions, EventBridge, SNS, S3, Glue, and Redshift**.

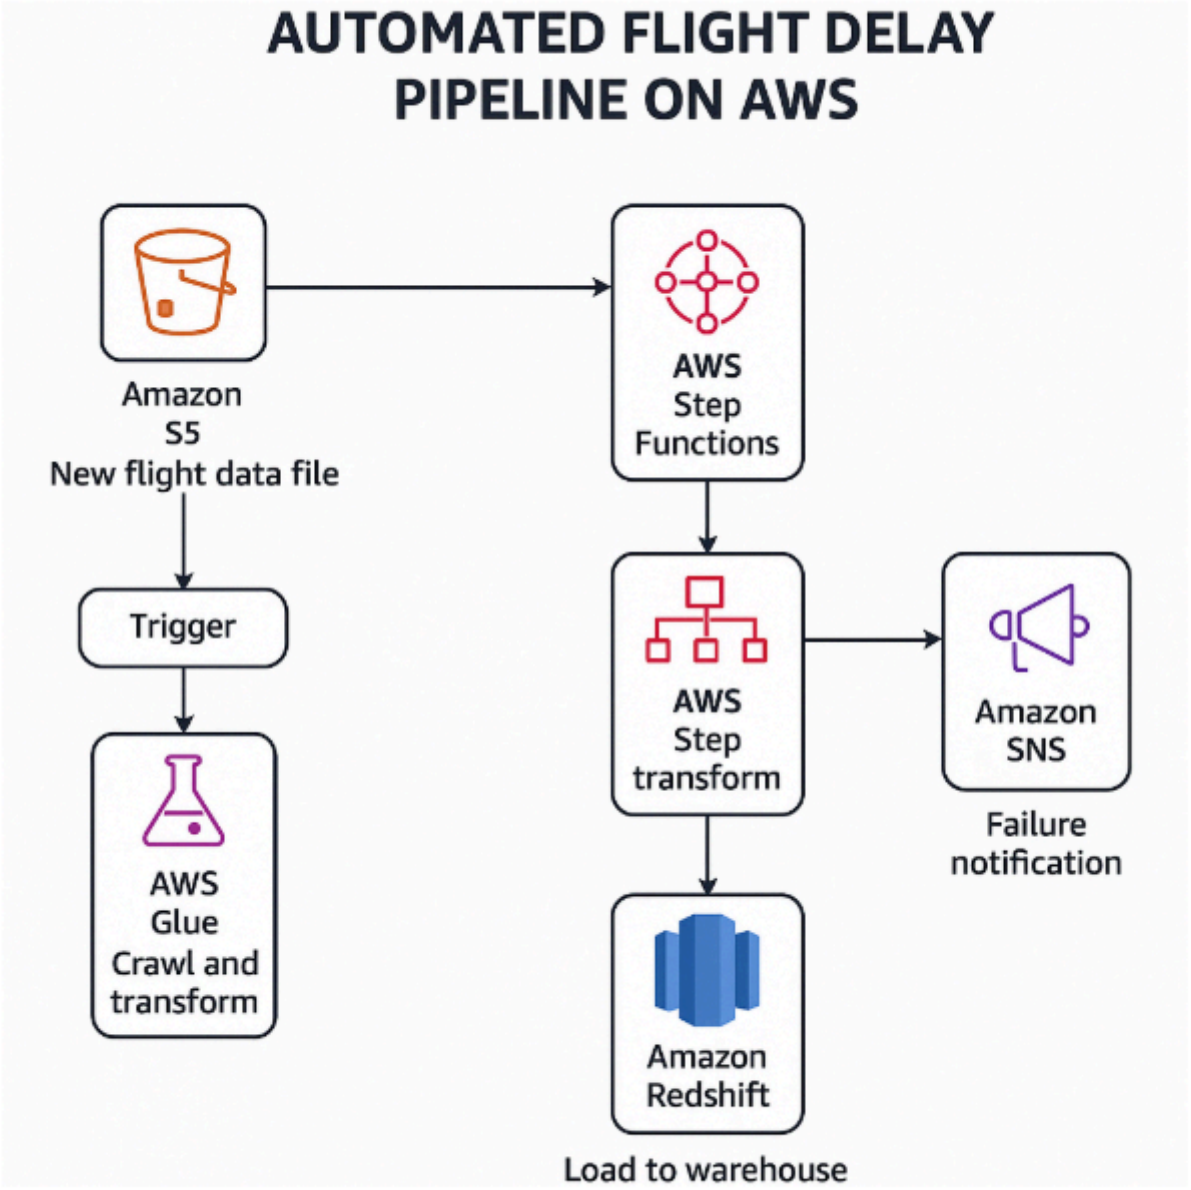

"Automated Flight Delay Processing Pipeline" project using **AWS Step Functions, EventBridge, SNS, S3, Glue, and Redshift**.

You can include this in your project README or technical documentation.

✈️ Automated Flight Delay Pipeline on AWS

📌 Project Overview

This project implements a **serverless, event-driven data pipeline** using AWS to automatically detect, transform, and analyze new flight data with a focus on **delays exceeding 60 minutes**. The solution ingests new data files added to an S3 bucket, processes and transforms them using Glue, and then loads the refined data into Amazon Redshift for analytics and reporting.

The pipeline also incorporates **Step Functions**, **EventBridge**, and **SNS** to automate the flow, monitor failures, and notify stakeholders in real time — **fully hands-free and scalable**.

🏗️ Architecture Summary

1. **Amazon S3** stores raw flight datasets.
2. **Amazon EventBridge** watches for new files added to the S3 bucket.
3. **AWS Step Functions** orchestrates the full data processing workflow.
4. **AWS Glue Crawler** scans and catalogs new data.
5. **AWS Glue Job** transforms the data to extract only the flights with delay times > 60 minutes.
6. **Amazon Redshift** stores the transformed data in a new analytics-ready table.
7. **Amazon SNS** sends failure notifications if any step fails, via email or other subscribers.

⚙️ Technologies Used

Service	Purpose
Amazon S3	Source bucket for raw flight data files
Amazon EventBridge	Detects new file uploads and triggers Step Functions
AWS Step Functions	Orchestrates data pipeline across Glue, Redshift, and SNS
AWS Glue Crawler	Automatically infers schema from new data
AWS Glue Job	Transforms the raw data (filters delay > 60 mins)
Amazon Redshift	Stores and visualizes the final results
Amazon SNS	Sends failure alerts (via email or SMS) on pipeline errors

---

🔁 Workflow Execution

1. **Trigger (EventBridge + S3):**
    - Whenever a new `.csv` or `.json` file is uploaded to the raw flight data S3 bucket, **EventBridge** detects this event.
    - It automatically **triggers the Step Function workflow** without any manual input.
2. **Crawling & Cataloging (Glue Crawler):**
    - A **Glue Crawler** scans the new dataset and updates the Data Catalog.
    - This ensures up-to-date schema metadata is available for transformation.
3. **Transformation (Glue Job):**
    - The **Glue Job** filters out flights with delay times greater than 60 minutes.
    - Data is cleaned, validated, and converted into a consistent format.
    - Output is stored in a separate S3 path or passed directly for loading into Redshift.
4. **Loading into Redshift:**
    - The filtered dataset is loaded into a **Redshift table called `flights_delayed_over_60`**.
    - This table supports further querying, dashboarding, or BI integrations.
5. **Error Notifications (SNS):**
    - If any stage fails (e.g., job error, Redshift load issue), the **Step Function triggers an SNS notification**.
    - A pre-configured email or SMS receives real-time alerts.

📊 Result

- A Redshift table named `flights_delayed_over_60` is populated with only those records where the delay exceeds 60 minutes.
- This enables quick downstream analytics to identify peak delay times, affected routes, or carrier patterns.

🚨 Error Handling & Monitoring

- **SNS Failures:** If Glue fails or Redshift load breaks, SNS sends a failure alert.
- **Step Function logs:** All executions (successful or failed) are logged in AWS Step Functions console.

- **CloudWatch Metrics:** Integrated monitoring available for each component, including Lambda logs (if used), job run time, crawler status, etc.

🔒 Security & Permissions

- IAM roles are tightly scoped for:
    - `StepFunctions-airline-data-pipeline-role` to run Glue, SNS, Redshift actions.
    - SNS topic is permissioned to only allow publishing from trusted roles.
    - S3 buckets have lifecycle rules and versioning enabled (optional).

🧪 Testing the Pipeline

To test the pipeline:

1. Upload a sample `.csv` or `.json` file with flight data to the S3 source bucket.
2. Watch the Step Function trigger and walk through each step.
3. Confirm the Redshift table is populated with correct delayed flights.
4. Introduce a failure (e.g., incorrect format) to verify SNS notification triggers.

🌟 Key Benefits

- **Fully automated**: No manual triggers required
- **Event-driven**: Responds in real-time to data arrival
- **Scalable**: Each component is serverless and cost-effective
- **Insight-ready**: Filtered data available in Redshift for instant BI consumption
- **Monitored**: End-to-end observability via logs, state tracing, and alerts

S3 bucket has been created. It has airports dim, date folder which will have updated data for everyday and a flights csv for loading data.

And going to create tables three tables for it one dimension and probably two facts When our trip need to create workgroups and name spaces as done in the previous projects and need to give access to the bucket

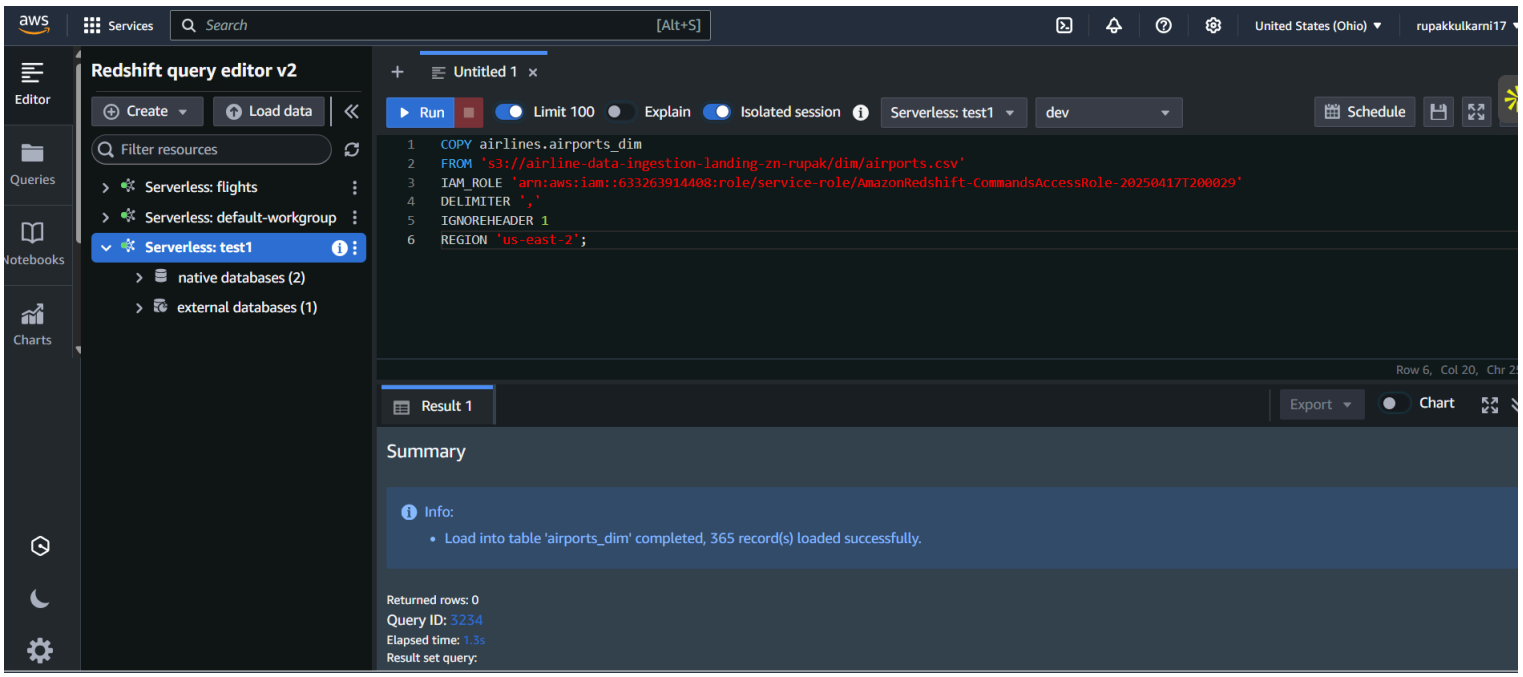

Opened inbound rules for this port 5439 used in redshift

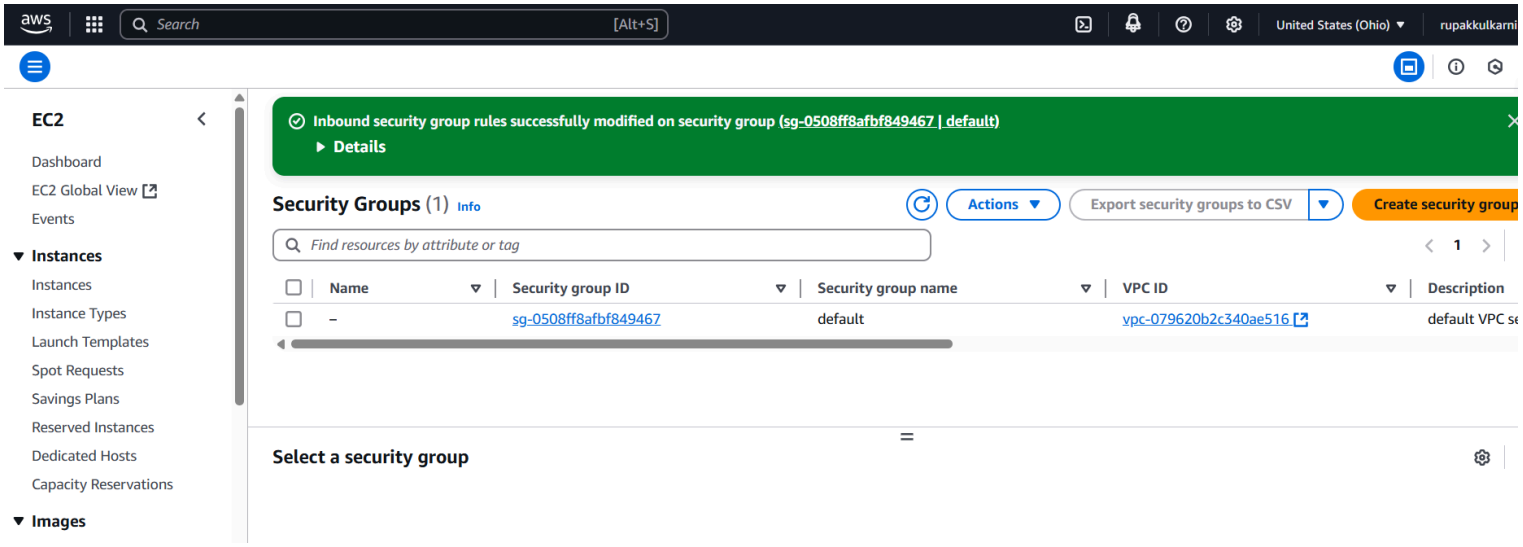

Glue connections:

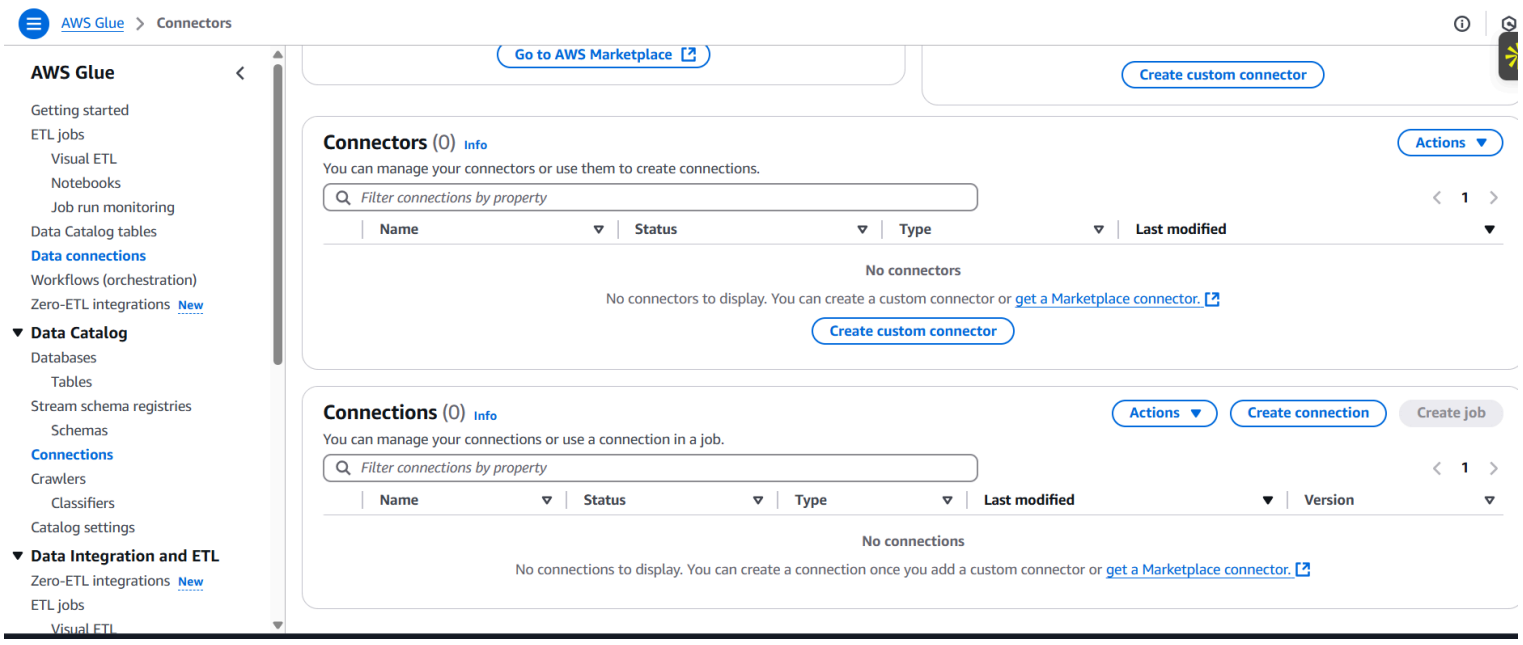

Connecting Glue with Redshift:

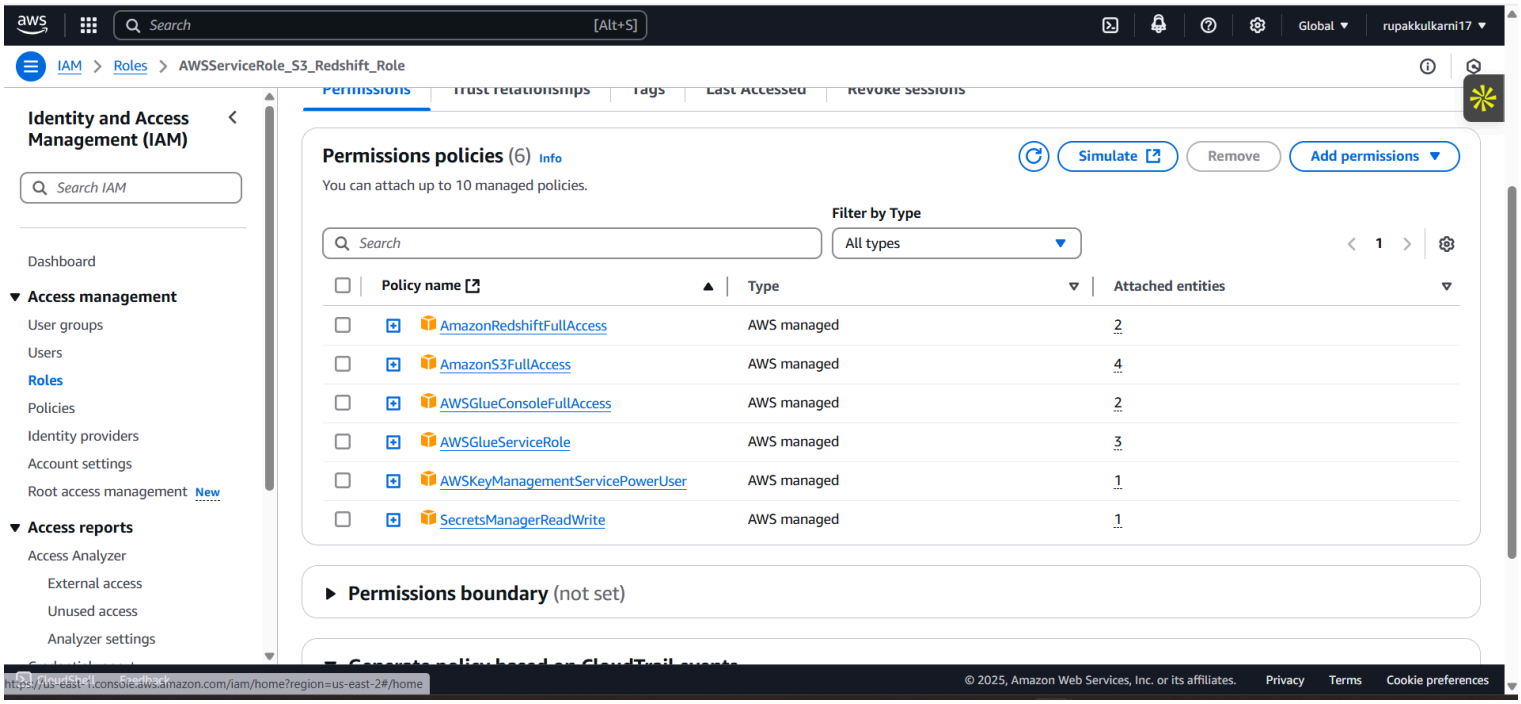

Inbound rules in redshift sandbox security group:

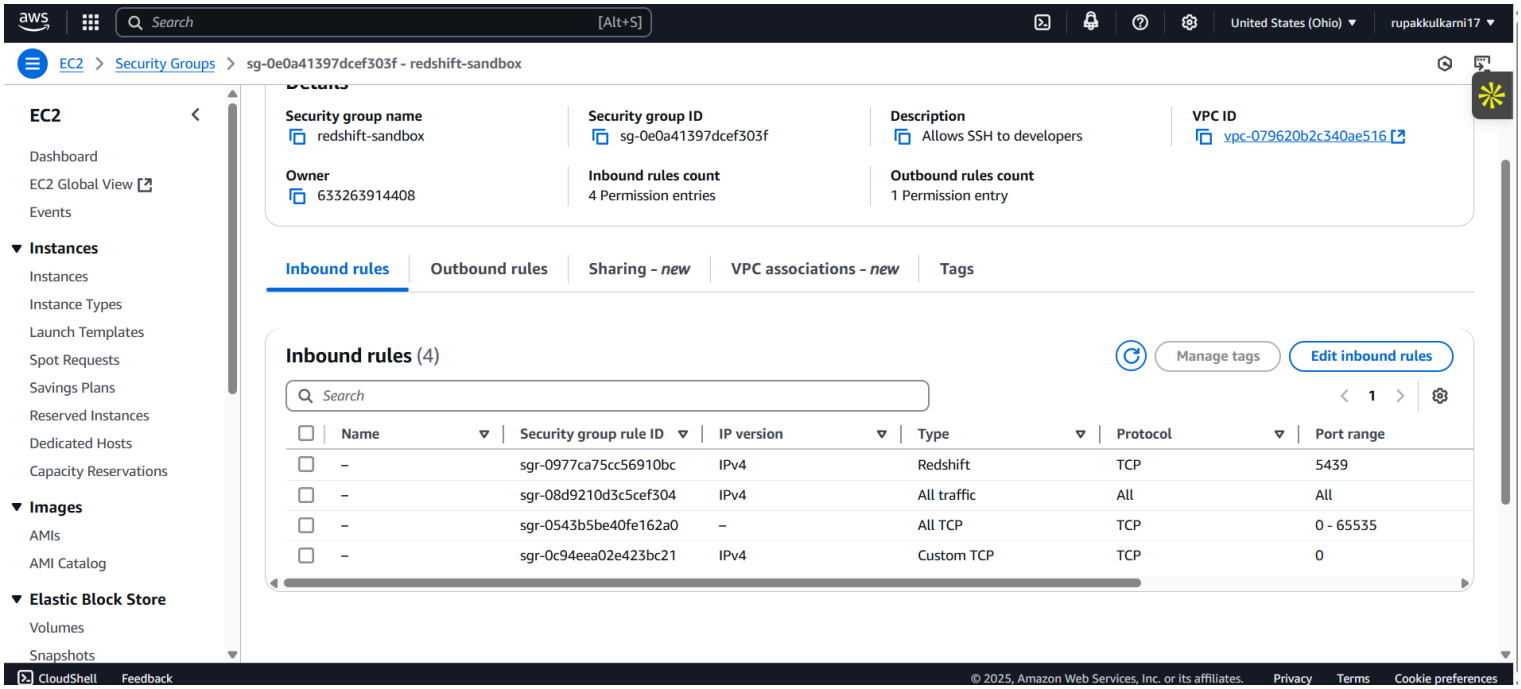

Creation of endpoints:

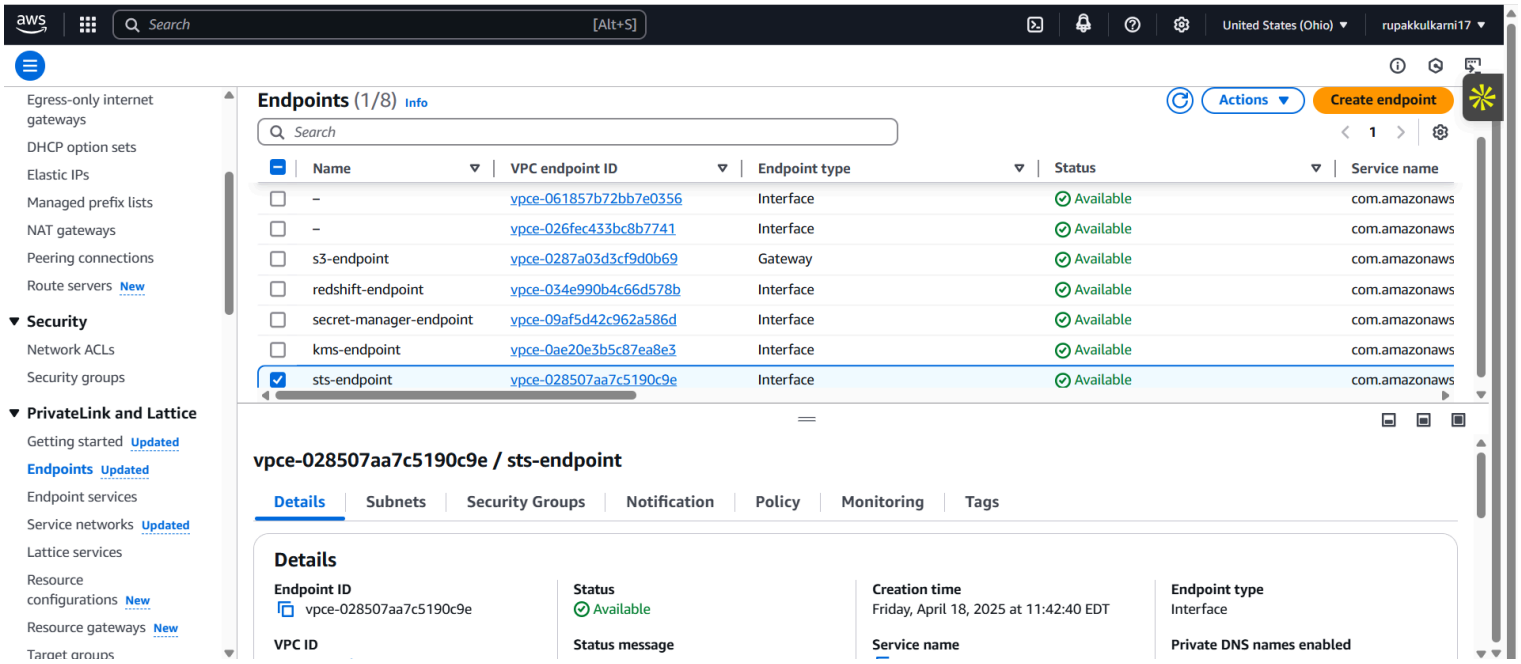

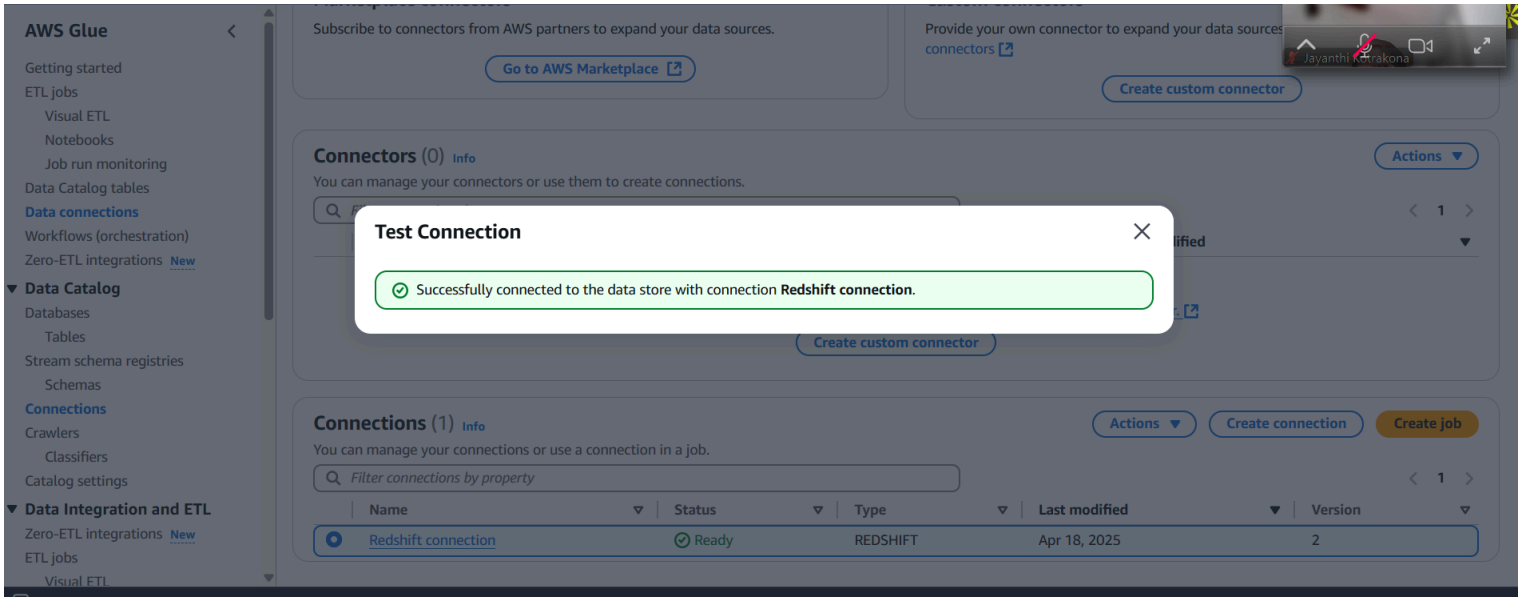

Data Transformation to push data in daily folder

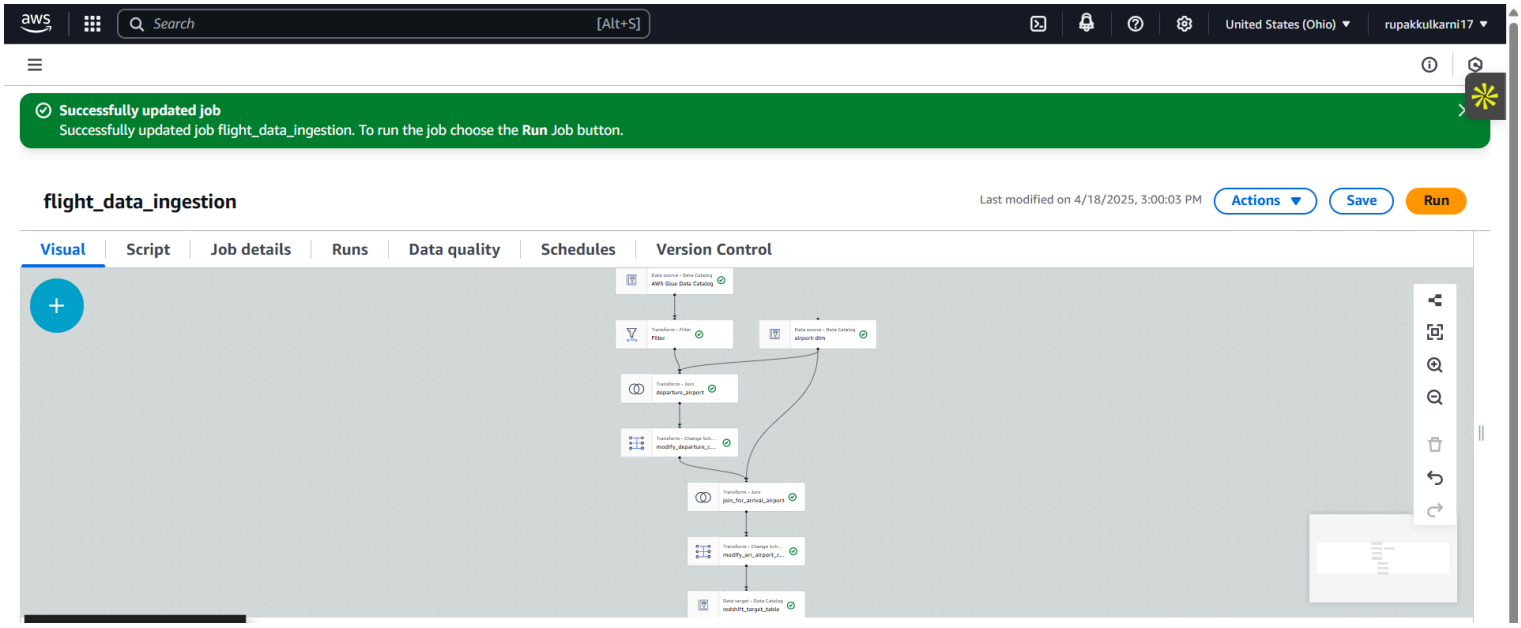

Turn on event bridge in the s3 bucket properties

Prepare step function:

Event bridge rule creation

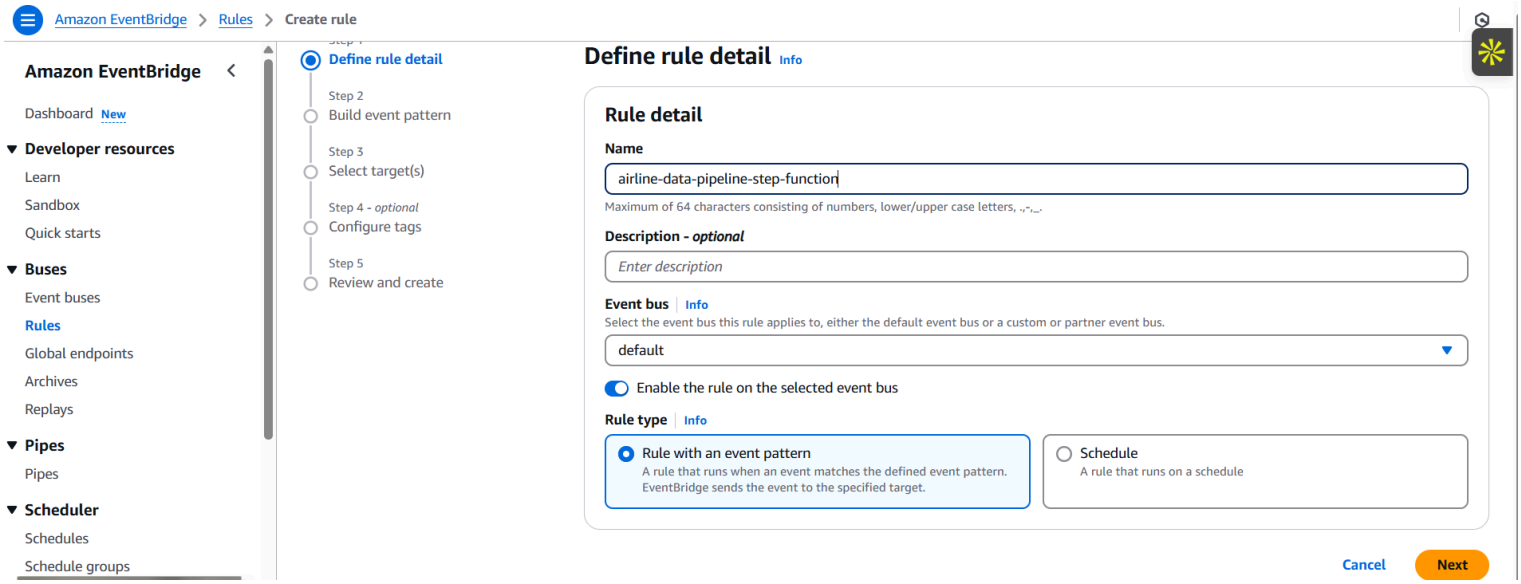

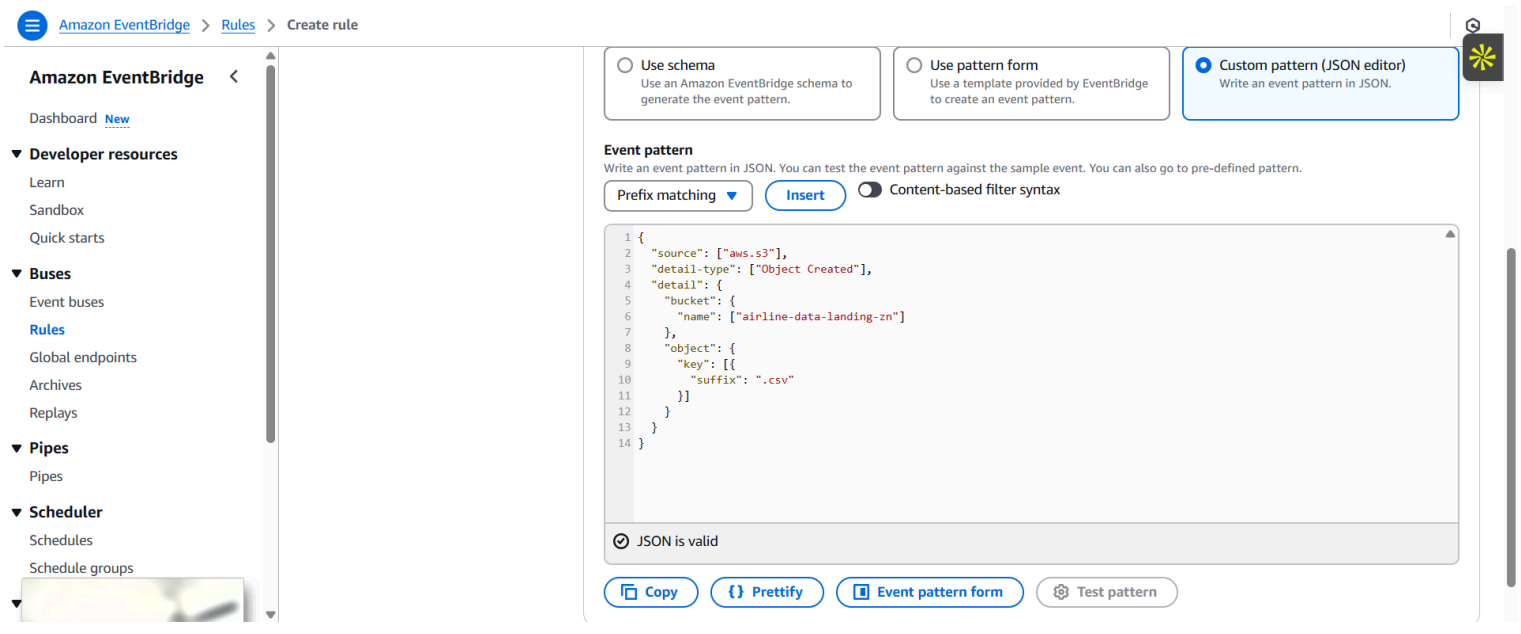

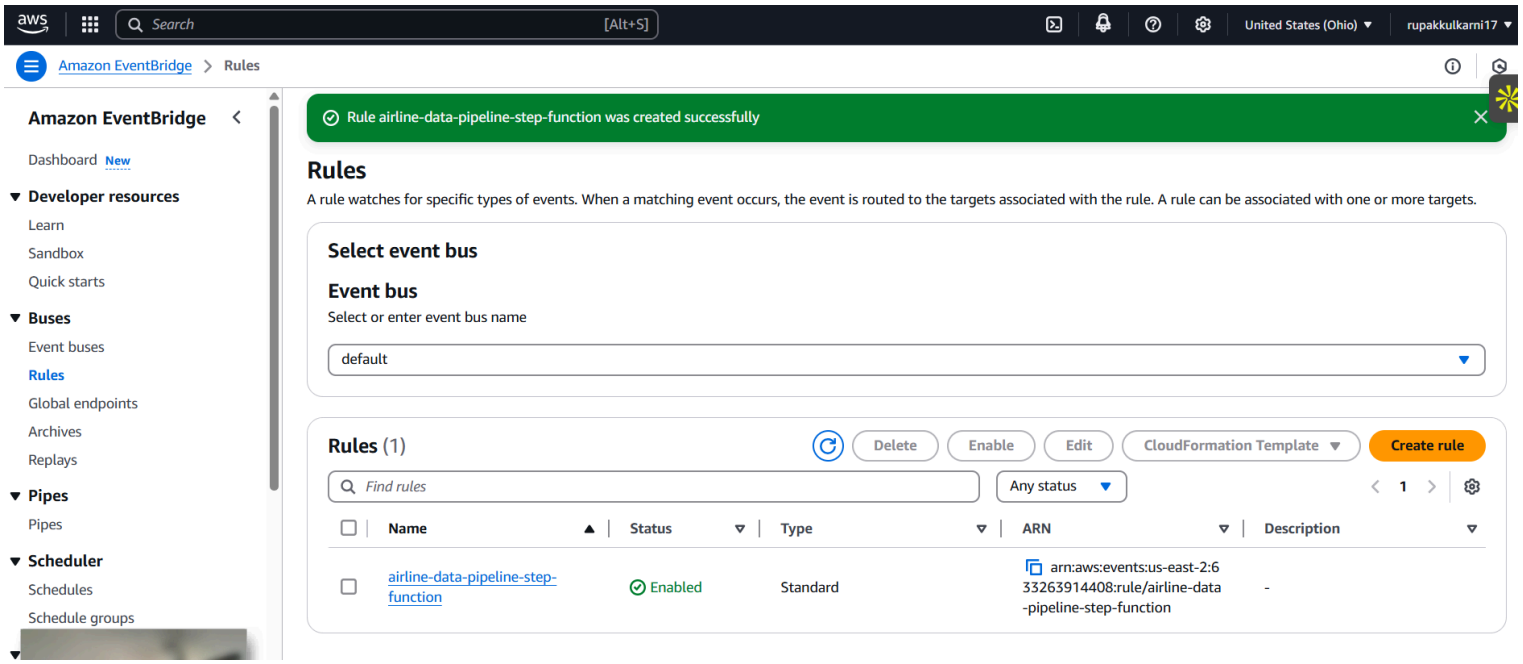

Now delete flight.csv. And again add it.

– you will see state machine is started automatically due to event bridge.
-state machine starts the crawler —-- all state run

▶ 2	TaskStateEntered	StartCrawler		00:00:00.025	Apr 18, 2025, 16:27:26.304 (UTC-04:00)
▶ 3	TaskScheduled	StartCrawler	aws-sdk:glue:startCrawler	00:00:00.025	Apr 18, 2025, 16:27:26.304 (UTC-04:00)
▶ 4	TaskStarted	StartCrawler	aws-sdk:glue:startCrawler	00:00:00.079	Apr 18, 2025, 16:27:26.358 (UTC-04:00)
▶ 5	TaskSucceeded	StartCrawler	aws-sdk:glue:startCrawler	00:00:00.511	Apr 18, 2025, 16:27:26.790 (UTC-04:00)
▶ 6	TaskStateExited	StartCrawler		00:00:00.530	Apr 18, 2025, 16:27:26.809 (UTC-04:00)
▶ 7	TaskStateEntered	GetCrawler		00:00:00.530	Apr 18, 2025, 16:27:26.809 (UTC-04:00)
▶ 8	TaskScheduled	GetCrawler	aws-sdk:glue:getCrawler	00:00:00.530	Apr 18, 2025, 16:27:26.809 (UTC-04:00)
▶ 9	TaskStarted	GetCrawler	aws-sdk:glue:getCrawler	00:00:00.592	Apr 18, 2025, 16:27:26.871 (UTC-04:00)
▶ 10	TaskSucceeded	GetCrawler	aws-sdk:glue:getCrawler	00:00:00.680	Apr 18, 2025, 16:27:26.959 (UTC-04:00)

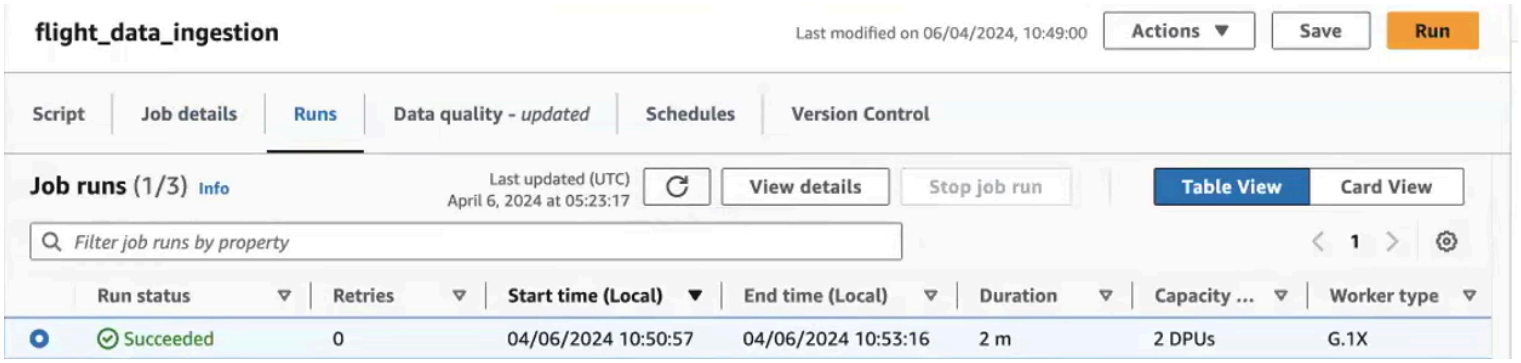

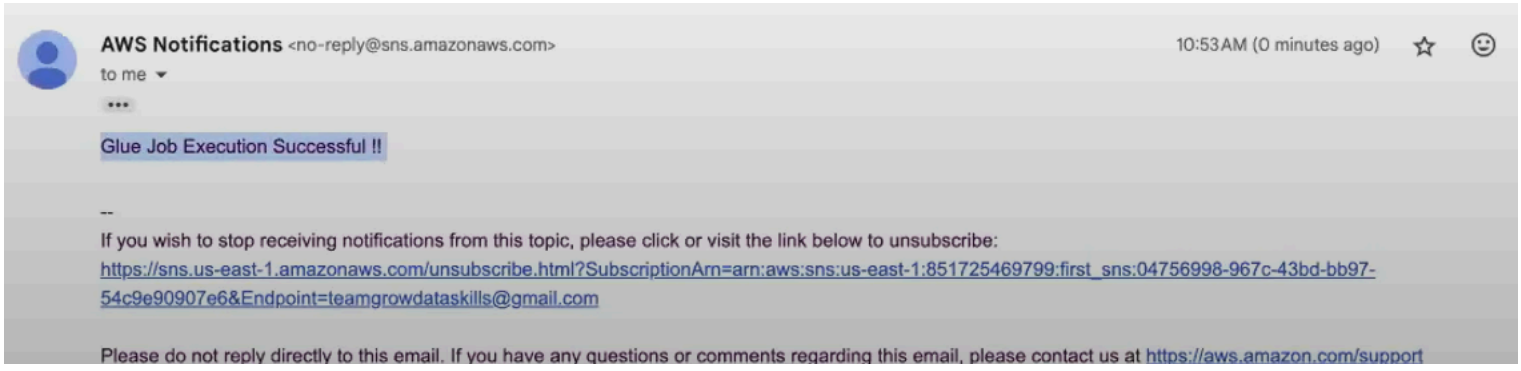

Departure delay greater than 60 minutes is seen in redshift cluster